

HIVE – HBASE CONNECTIVITY

Step -1 : Create a Hive table ,which acts as stage table to hold data for operation

```
hive > CREATE TABLE sales_order_csv
(
  order_number int,
  qty_ordered int,
  price_each int,
  order_line_number int,
  sales int,
  status STRING,
  qtr_id int,
  month_id int,
  year_id int,
  product_line STRING,
  msrp int,
  product_code STRING,
  phone string,
  city STRING,
  state string,
  postal_code STRING,
  country STRING,
  territory STRING,
  contact_lname STRING,
  contact_fname STRING,
  dealsize STRING
)
ROW FORMAT DELIMITED
FIELDS TERMINATED BY ','
TBLPROPERTIES("skip.header.line.count"="1");
```

```
>
>
> ROW FORMAT DELIMITED
> FIELDS TERMINATED BY ','
> TBLPROPERTIES("skip.header.line.count"="1");
OK
Time taken: 2.482 seconds
hive> show tables
> ;
OK
sales_order_csv
Time taken: 0.219 seconds, Fetched: 1 row(s)
```

Step -2 : Load data from csv file to stage table created in Step-1

\$hive > load data local inpath '/tmp/sales_order_data.csv' into table sales_order_csv;

```
hive> load data local inpath '/tmp/sales_order_data.csv' into table sales_order_csv;
Loading data to table default.sales_order_csv
Table default.sales_order_csv stats: [numFiles=1, totalSize=360233]
OK
Time taken: 0.988 seconds
```

Step -3 : Create Hbase table with column families

(Order number will act as row_key, hence it is not considered in column family while creating table)

hbase(main):013:0> create

'sales_order_data_hbase','qty_ordered','price_each','order_line_number','sales','status','qtr_id','month_id','year_id','product_line','msrp','product_code','phone','city','state','postal_code','country','territory','contact_lname','contact_fname','dealsize'

Step -4 : Create Hive table with Hbase SerDe properties

```
CREATE EXTERNAL TABLE IF NOT EXISTS sales_order_hive_hbase_connect1
(
key int,
qty_ordered int,
price_each int,
order_line_number int,
sales int,
status STRING,
qtr_id int,
month_id int,
year_id int,
product_line STRING,
msrp int,
product_code STRING,
phone string,
city STRING,
state string,
postal_code STRING,
country STRING,
territory STRING,
```

```
contact_lname STRING,  
contact_fname STRING,  
dealsize STRING  
)  
ROW FORMAT DELIMITED  
FIELDS TERMINATED BY ','  
STORED BY 'org.apache.hadoop.hive.hbase.HBaseStorageHandler'  
WITH SERDEPROPERTIES  
("hbase.columns.mapping"=":key,  
qty_ordered:qty_ordered,  
price_each:price_each,  
order_line_number:order_line_number,  
sales:sales,  
status:status,  
qtr_id:qtr_id,  
month_id:month_id,  
year_id:year_id,  
product_line:product_line,  
msrp:msrp,  
product_code:product_code,  
phone:phone,  
city:city,  
state:state,  
postal_code:postal_code,  
country:country,  
territory:territory,  
contact_lname:contact_lname,  
contact_fname:contact_fname,  
dealsize:dealsize")  
TBLPROPERTIES("hbase.table.name"="sales_order_data_hbase");
```

Step -5 : Insert data from stage table to Hive table with SerDe properties

hive > Insert into table sales_order_hive_hbase_connect select * from sales_order_csv;

```
hive> Insert into table sales_order_hive_hbase_connect1 select * from sales_order_csv
> ;
Query ID = ccloudera_20221128082727_6df2f5eb-2b73-42d5-a6b1-ff278f5cf397
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks is set to 0 since there's no reduce operator
Starting Job = job_1669648232524_0001, Tracking URL = http://quickstart.cloudera:8088/proxy/application_1669648232524_0001/
Kill Command = /usr/lib/hadoop/bin/hadoop job -kill job_1669648232524_0001
Hadoop job information for Stage-0: number of mappers: 1; number of reducers: 0
2022-11-28 08:28:46,545 Stage-0 map = 0%, reduce = 0%
2022-11-28 08:29:23,150 Stage-0 map = 100%, reduce = 0%, Cumulative CPU 6.27 sec
MapReduce Total cumulative CPU time: 6 seconds 270 msec
Ended Job = job_1669648232524_0001
MapReduce Jobs Launched:
Stage-Stage-0: Map: 1 Cumulative CPU: 6.27 sec HDFS Read: 375245 HDFS Write: 0 SUCCESS
Total MapReduce CPU Time Spent: 6 seconds 270 msec
OK
Time taken: 89.353 seconds
hive>
```

Step -6 : Inspecting data

1). Hive table

Hive > Select * from sales_order_hive_hbase_connect limit 10;

```
hive> select * from sales_order_hive_hbase_connect1 limit 10;
OK
10100 49 34 1 1689 Shipped 1 1 2003 Vintage Cars 41 S24_3969 6035558647 Nashua NH 62005 USA N
A Young Valarie Small
10101 46 53 2 2472 Shipped 1 1 2003 Vintage Cars 44 S24_2022 +49 69 66 90 2555 Frankfurt 6
0528 Germany EMEA Keitel Roland Small
10102 41 50 1 2055 Shipped 1 1 2003 Vintage Cars 53 S18_1367 2125551500 NYC NY 10022 USA N
A Frick Michael Small
10103 42 100 6 4460 Shipped 1 1 2003 Classic Cars 101 S700_2824 07-98 9555 Stavern 4110 Norwa
y EMEA Bergulfsen Jonas Medium
10104 32 53 2 1705 Shipped 1 1 2003 Trains 58 S50_1514 (91) 555 94 44 Madrid 28034 Spain EMEAF
reyre Diego Small
10105 25 56 8 1419 Shipped 1 2 2003 Ships 54 S72_3212 31 12 3555 Kobenhavn 1734 Denma
rk EMEA Petersen Jytte Small
10106 48 52 15 2526 Shipped 1 2 2003 Planes 49 S72_1253 035-640555 Bergamo 24100 Italy EMEAR
owelll Giovanni Small
10107 20 92 8 1858 Shipped 1 2 2003 Motorcycles 99 S32_1374 2125557818 NYC NY 10022 USA N
A Yu Kwai Small
10108 34 82 14 2821 Shipped 1 3 2003 Motorcycles 81 S50_4713 +63 2 555 3587 Makati City 1227
WM Philippines Japan Cruz Arnold Small
10109 29 32 6 930 Shipped 1 3 2003 Classic Cars 37 S24_2972 2155559857 Philadelphia PA 71270
USA NA Hernandez Rosa Small
Time taken: 0.189 seconds, Fetched: 10 row(s)
hive>
```

2). Hbase table – before loading data

Hbase > scan 'sales_order_data_hbase'

```
=> Hbase::Table - sales_order_data_hbase
hbase(main):019:0> scan 'sales_order_data_hbase'
ROW COLUMN+CELL
0 row(s) in 0.2580 seconds
hbase(main):020:0>
```

3). Hbase table – After loading data

Hbase > scan 'sales_order_data_hbase'

```
10425      column=phone:phone, timestamp=1669652961409, value=40.67.8555
10425      column=postal_code:postal_code, timestamp=1669652961409, value=44000
10425      column=price_each:price_each, timestamp=1669652961409, value=100
10425      column=product_code:product_code, timestamp=1669652961409, value=S50_1392
10425      column=product_line:product_line, timestamp=1669652961409, value=Trucks and Buses
10425      column=qtr_id:qtr_id, timestamp=1669652961409, value=2
10425      column=qty_ordered:qty_ordered, timestamp=1669652961409, value=18
10425      column=sales:sales, timestamp=1669652961409, value=1895
10425      column=state:state, timestamp=1669652961409, value=
10425      column=status:status, timestamp=1669652961409, value=In Process
10425      column=territory:territory, timestamp=1669652961409, value=EMEA
10425      column=year_id:year_id, timestamp=1669652961409, value=2005
307 row(s) in 9.1850 seconds
hbase(main):021:0>
```

4). Loading sample data in Hbase table

Hbase > put 'sales_order_data_hbase','11111','qty_ordered: qty_ordered','10'

```
0 row(s) in 0.3350 seconds

hbase(main):044:0> put 'sales_order_data_hbase','11111','qty_ordered:qty_ordered','10'
0 row(s) in 0.0060 seconds

hbase(main):045:0>
```

5). Accessing hbase table data from Hive

Hive > Select * from sales_order_hive_hbase_connect where key = 11111;

```
hive> select * from sales_order_hive_hbase_connect1 where key = 11111;
OR
11111  10      NULL  NULL  NULL  NULL  NULL  NULL  NULL  NULL  NULL  NULL  NULL  NULL  NULL  NULL  NULL  NULL
NULL
Time taken: 0.362 seconds, Fetched: 1 row(s)
hive>
```